

I] Importing Dataset and Basic Check

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')
```

1) Importing Dataset

```
path=r'C:\Users\venky\Downloads\ecommerce_recommendation_dataset.csv'
df=pd.read_csv(path)
df.head()
```

	user_id	product_id	category	price	rating	review_count
0	78517	1645	Books	842.23	2	155
1	52887	100	Books	253.76	3	331
2	59395	585	Books	483.65	2	236
3	54739	3774	Groceries	459.37	2	227
4	42723	2119	Groceries	150.11	2	214

	user_gender	user_location	purchase_history	...
0	Other	Urban	False	...
1	Other	Suburban	False	...
2	Female	Rural	True	...
3	Female	Urban	False	...
4	Female	Urban	True	...

	review_sentiment_score	user_engagement_score	ad_click_rate
0	-0.28	0.68	0.04
1	0.28	0.11	0.89

2		0.23		0.35		0.99
Evening						
3		0.93		0.73		0.16
Afternoon						
4		0.11		0.26		0.17
Night						

	day_of_week	season	payment_method	coupon_used	product_popularity
0	Thursday	Summer	Debit Card	False	0.54
1	Saturday	Summer	Debit Card	False	0.77
2	Tuesday	Fall	Debit Card	False	0.14
3	Tuesday	Spring	Credit Card	False	0.18
4	Wednesday	Spring	PayPal	False	0.66

[5 rows x 51 columns]

2) Checking Shape and Size of Dataset

```
print(f"shape of data set is = {df.shape}")
print(f"size of data set is = {df.size}")

shape of data set is = (60000, 51)
size of data set is = 3060000
```

3) Checking Duplicate Rows and Missing Value

```
print(f"total duplicate rows in dataset is = {df.duplicated().sum()}")
print()
print(f"total missing values rows in dataset is = {df.isnull().sum().sum()}")
print()
print(f"missing data in columns is as below\
n{df.isnull().sum()}", sep='')

total duplicate rows in dataset is = 0

total missing values rows in dataset is = 0

missing data in columns is as below
user_id          0
product_id       0
category         0
price           0
rating          0
```

review_count	0
user_age	0
user_gender	0
user_location	0
purchase_history	0
time_on_page	0
add_to_cart_count	0
search_keywords	0
discount_applied	0
user_membership	0
user_browser	0
user_device	0
purchase_time	0
session_duration	0
clicks_on_ads	0
page_views	0
referral_source	0
wishlist_additions	0
cart_abandonment_rate	0
average_spent	0
user_income	0
user_education	0
user_marital_status	0
product_availability	0
stock_status	0
product_return_rate	0
product_color	0
product_size	0
is_top_seller	0
discount_percentage	0
time_to_purchase	0
delivery_time	0
shipping_fee	0
seller_rating	0
seller_response_time	0
seller_location	0
product_rating_variance	0
review_sentiment_score	0
user_engagement_score	0
ad_click_rate	0
time_of_day	0
day_of_week	0
season	0
payment_method	0
coupon_used	0
product_popularity	0
dtype: int64	

df.info()

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 60000 entries, 0 to 59999
```

```
Data columns (total 51 columns):
```

#	Column	Non-Null Count	Dtype
0	user_id	60000 non-null	int64
1	product_id	60000 non-null	int64
2	category	60000 non-null	object
3	price	60000 non-null	float64
4	rating	60000 non-null	int64
5	review_count	60000 non-null	int64
6	user_age	60000 non-null	int64
7	user_gender	60000 non-null	object
8	user_location	60000 non-null	object
9	purchase_history	60000 non-null	bool
10	time_on_page	60000 non-null	float64
11	add_to_cart_count	60000 non-null	int64
12	search_keywords	60000 non-null	object
13	discount_applied	60000 non-null	bool
14	user_membership	60000 non-null	object
15	user_browser	60000 non-null	object
16	user_device	60000 non-null	object
17	purchase_time	60000 non-null	object
18	session_duration	60000 non-null	float64
19	clicks_on_ads	60000 non-null	int64
20	page_views	60000 non-null	int64
21	referral_source	60000 non-null	object
22	wishlist_additions	60000 non-null	int64
23	cart_abandonment_rate	60000 non-null	float64
24	average_spent	60000 non-null	float64
25	user_income	60000 non-null	float64
26	user_education	60000 non-null	object
27	user_marital_status	60000 non-null	object
28	product_availability	60000 non-null	object
29	stock_status	60000 non-null	object
30	product_return_rate	60000 non-null	float64
31	product_color	60000 non-null	object
32	product_size	60000 non-null	object
33	is_top_seller	60000 non-null	bool
34	discount_percentage	60000 non-null	float64
35	time_to_purchase	60000 non-null	float64
36	delivery_time	60000 non-null	float64
37	shipping_fee	60000 non-null	float64
38	seller_rating	60000 non-null	int64
39	seller_response_time	60000 non-null	float64
40	seller_location	60000 non-null	object
41	product_rating_variance	60000 non-null	float64
42	review_sentiment_score	60000 non-null	float64
43	user_engagement_score	60000 non-null	float64
44	ad_click_rate	60000 non-null	float64

```

45  time_of_day          60000 non-null object
46  day_of_week          60000 non-null object
47  season               60000 non-null object
48  payment_method       60000 non-null object
49  coupon_used          60000 non-null bool
50  product_popularity   60000 non-null float64
dtypes: bool(4), float64(17), int64(10), object(20)
memory usage: 21.7+ MB

```

4) Checking DataTypes

```
df.dtypes
```

```

user_id          int64
product_id       int64
category         object
price            float64
rating           int64
review_count     int64
user_age         int64
user_gender      object
user_location    object
purchase_history bool
time_on_page     float64
add_to_cart_count int64
search_keywords  object
discount_applied bool
user_membership  object
user_browser     object
user_device      object
purchase_time    object
session_duration float64
clicks_on_ads    int64
page_views       int64
referral_source  object
wishlist_additions int64
cart_abandonment_rate float64
average_spent    float64
user_income      float64
user_education   object
user_marital_status object
product_availability object
stock_status     object
product_return_rate float64
product_color    object
product_size     object
is_top_seller    bool
discount_percentage float64
time_to_purchase float64

```

```
delivery_time          float64
shipping_fee           float64
seller_rating          int64
seller_response_time   float64
seller_location        object
product_rating_variance float64
review_sentiment_score float64
user_engagement_score  float64
ad_click_rate          float64
time_of_day            object
day_of_week            object
season                 object
payment_method         object
coupon_used            bool
product_popularity     float64
dtype: object
```

5) Checking Target Column - 'purchase_history'

```
df['purchase_history'].value_counts(normalize=True)*100

purchase_history
True      50.146667
False     49.853333
Name: proportion, dtype: float64
```

6) Feature Engineering of 'purchase_time'

```
print("purchase_time")
print("number of unique values :",df['purchase_time'].nunique())
print("max :",df['purchase_time'].max())
print("min :",df['purchase_time'].min())
print("datatype :",df['purchase_time'].dtype)

purchase_time
number of unique values : 8749
max : 2024-12-31 00:00:00
min : 2024-01-01 00:00:00
datatype : object
```

6.a) Converting To datetime datatype

```
df['purchase_time']=pd.to_datetime(df['purchase_time'])
print("datatype :",df['purchase_time'].dtype)

datatype : datetime64[ns]
```

6.b) Extracting values from date time

```
print("max year = ",df['purchase_time'].dt.year.max())
print("min year = ",df['purchase_time'].dt.year.min())
print()
print("max month = ",df['purchase_time'].dt.month.max())
print("min month = ",df['purchase_time'].dt.month.min())
# df['day'] = df['purchase_time'].dt.day
# df['month'] = df['purchase_time'].dt.month
# df['year'] = df['purchase_time'].dt.year
# df['hour'] = df['purchase_time'].dt.hour
# df['minute'] = df['purchase_time'].dt.minute

max year = 2024
min year = 2024

max month = 12
min month = 1

df['purchase_time_month']=df['purchase_time'].dt.month
df['purchase_time_month'].value_counts()

purchase_time_month
5      5174
12     5163
10     5152
3       5105
1       5044
7       5010
8       4985
4       4982
6       4920
11      4884
2       4801
9       4780
Name: count, dtype: int64

'time_of_day'
print("time_of_day")
print("number of unique values :",df['time_of_day'].nunique())
print("value counts are :\n",df['time_of_day'].value_counts(),sep='')
print("datatype :",df['time_of_day'].dtype)

time_of_day
number of unique values : 4
value counts are :
time_of_day
Night      15095
Evening    15025
Morning    14968
Afternoon  14912
```

Name: count, dtype: int64

datatype : object

```
df.drop(columns=['purchase_time'],axis=1,inplace=True)
```

```
basic_info=pd.DataFrame()
```

```
basic_info['feature_name']=df.columns
```

```
basic_info['missing_values']=df.isnull().sum().values
```

```
basic_info['data_types']=df.dtypes.values
```

```
basic_info['number_of_uniquevalues']=df.nunique().values
```

```
basic_info
```

	feature_name	missing_values	data_types
number_of_uniquevalues			
0	user_id	0	int64
45154			
1	product_id	0	int64
4999			
2	category	0	object
5			
3	price	0	float64
45014			
4	rating	0	int64
5			
5	review_count	0	int64
500			
6	user_age	0	int64
52			
7	user_gender	0	object
3			
8	user_location	0	object
3			
9	purchase_history	0	bool
2			
10	time_on_page	0	float64
2951			
11	add_to_cart_count	0	int64
10			
12	search_keywords	0	object
5			
13	discount_applied	0	bool
2			
14	user_membership	0	object
4			
15	user_browser	0	object
4			
16	user_device	0	object
3			
17	session_duration	0	float64
55235			

18	clicks_on_ads	0	int64
20			
19	page_views	0	int64
99			
20	referral_source	0	object
4			
21	wishlist_additions	0	int64
20			
22	cart_abandonment_rate	0	float64
101			
23	average_spent	0	float64
56634			
24	user_income	0	float64
59912			
25	user_education	0	object
4			
26	user_marital_status	0	object
4			
27	product_availability	0	object
3			
28	stock_status	0	object
3			
29	product_return_rate	0	float64
101			
30	product_color	0	object
5			
31	product_size	0	object
3			
32	is_top_seller	0	bool
2			
33	discount_percentage	0	float64
5001			
34	time_to_purchase	0	float64
25810			
35	delivery_time	0	float64
1401			
36	shipping_fee	0	float64
5001			
37	seller_rating	0	int64
5			
38	seller_response_time	0	float64
7100			
39	seller_location	0	object
3			
40	product_rating_variance	0	float64
201			
41	review_sentiment_score	0	float64
201			
42	user_engagement_score	0	float64

```

101
43          ad_click_rate          0    float64
101
44          time_of_day            0    object
4
45          day_of_week            0    object
7
46          season                 0    object
4
47          payment_method         0    object
4
48          coupon_used            0    bool
2
49          product_popularity     0    float64
101
50          purchase_time_month    0    int32
12

y=df['purchase_history']
x=df.drop(columns=['purchase_history'],axis=1)

categorical=[]
numerical=[]
descrete=[]
x.drop(columns=['user_id', 'product_id'],axis=1,inplace=True)
for i in x.columns:
    if x[i].nunique()>25 and i not in ['user_id','product_id']:
        numerical.append(i)
    elif x[i].nunique()<25 and x[i].dtype=='object':
        categorical.append(i)
    else:
        descrete.append(i)

print(numerical)
print()
print("total numerical features are =",len(numerical))

['price', 'review_count', 'user_age', 'time_on_page',
'session_duration', 'page_views', 'cart_abandonment_rate',
'average_spent', 'user_income', 'product_return_rate',
'discount_percentage', 'time_to_purchase', 'delivery_time',
'shipping_fee', 'seller_response_time', 'product_rating_variance',
'review_sentiment_score', 'user_engagement_score', 'ad_click_rate',
'product_popularity']

total numerical features are = 20

print(categorical)
print()
print("total categorical features are =",len(categorical))

```

```
['category', 'user_gender', 'user_location', 'search_keywords',
'user_membership', 'user_browser', 'user_device', 'referral_source',
'user_education', 'user_marital_status', 'product_availability',
'stock_status', 'product_color', 'product_size', 'seller_location',
'time_of_day', 'day_of_week', 'season', 'payment_method']

total categorical features are = 19

print(descrete)
print()
print("total descrete features are =",len(descrete))

['rating', 'add_to_cart_count', 'discount_applied', 'clicks_on_ads',
'wishlist_additions', 'is_top_seller', 'seller_rating', 'coupon_used',
'purchase_time_month']

total descrete features are = 9
```

II] Feature Selection & Multi-Colinearity

1) Multi-Colinearity in Numerical Features

1.a) Corellation Matrix

```
corrmatrix=df[numerical].corr()
corrmatrix
```

	price	review_count	user_age	
time_on_page \				
price	1.000000	0.000553	-0.005285	-
0.002197				
review_count	0.000553	1.000000	-0.000759	
0.002428				
user_age	-0.005285	-0.000759	1.000000	
0.004028				
time_on_page	-0.002197	0.002428	0.004028	
1.000000				
session_duration	-0.000339	-0.000119	0.000812	-
0.001737				
page_views	0.002720	-0.001148	-0.002092	
0.002644				
cart_abandonment_rate	0.004704	-0.004129	-0.001691	
0.002740				
average_spent	-0.002186	0.001596	-0.002669	-
0.000698				
user_income	-0.002944	-0.001625	0.001072	-
0.001668				
product_return_rate	0.002338	0.004427	0.002673	-

0.004414				
discount_percentage	0.006372	0.000887	-0.003565	
0.002988				
time_to_purchase	-0.004822	0.001550	-0.004478	
0.003364				
delivery_time	0.000883	0.001659	-0.000821	-
0.002803				
shipping_fee	0.001692	0.003383	-0.000638	
0.009305				
seller_response_time	0.000019	-0.000159	-0.003831	-
0.000462				
product_rating_variance	0.011372	-0.004260	0.002916	
0.003512				
review_sentiment_score	-0.009115	0.002156	-0.003348	
0.000330				
user_engagement_score	-0.000325	0.006545	-0.001121	-
0.000154				
ad_click_rate	0.004739	0.005577	0.001147	-
0.000799				
product_popularity	-0.001021	0.004671	0.002571	-
0.004101				
	session_duration	page_views		
cart_abandonment_rate \				
price	-0.000339	0.002720		
0.004704				
review_count	-0.000119	-0.001148		-
0.004129				
user_age	0.000812	-0.002092		-
0.001691				
time_on_page	-0.001737	0.002644		
0.002740				
session_duration	1.000000	-0.005290		-
0.003657				
page_views	-0.005290	1.000000		
0.005704				
cart_abandonment_rate	-0.003657	0.005704		
1.000000				
average_spent	0.003399	0.005138		-
0.010053				
user_income	0.001579	0.002168		
0.003268				
product_return_rate	0.003760	0.000223		
0.001191				
discount_percentage	0.002069	0.003101		-
0.000313				
time_to_purchase	-0.000845	-0.003670		
0.004395				
delivery_time	0.008153	0.002657		-

0.005104			
shipping_fee	-0.000676	0.002222	
0.001536			
seller_response_time	0.007093	0.000090	
0.003947			
product_rating_variance	-0.000166	0.000418	
0.007186			
review_sentiment_score	0.001741	-0.002711	-
0.003193			
user_engagement_score	-0.000691	-0.003882	
0.003944			
ad_click_rate	0.002220	0.003758	
0.007245			
product_popularity	-0.003204	-0.005096	
0.005538			

	average_spent	user_income	
product_return_rate \			
price	-0.002186	-0.002944	
0.002338			
review_count	0.001596	-0.001625	
0.004427			
user_age	-0.002669	0.001072	
0.002673			
time_on_page	-0.000698	-0.001668	-
0.004414			
session_duration	0.003399	0.001579	
0.003760			
page_views	0.005138	0.002168	
0.000223			
cart_abandonment_rate	-0.010053	0.003268	
0.001191			
average_spent	1.000000	-0.002158	
0.003798			
user_income	-0.002158	1.000000	-
0.004187			
product_return_rate	0.003798	-0.004187	
1.000000			
discount_percentage	0.008415	0.000450	-
0.001619			
time_to_purchase	-0.004868	0.000561	
0.002045			
delivery_time	0.009656	-0.000427	
0.001719			
shipping_fee	0.003372	0.003680	
0.004030			
seller_response_time	0.000699	-0.003094	
0.001394			
product_rating_variance	-0.002859	0.001301	

0.003918			
review_sentiment_score	0.006323	0.005753	
0.003248			
user_engagement_score	-0.000966	-0.000097	
0.001266			
ad_click_rate	-0.005269	-0.004528	
0.001472			
product_popularity	0.004336	-0.002315	-
0.002570			
	discount_percentage	time_to_purchase	
delivery_time \			
price	0.006372	-0.004822	
0.000883			
review_count	0.000887	0.001550	
0.001659			
user_age	-0.003565	-0.004478	-
0.000821			
time_on_page	0.002988	0.003364	-
0.002803			
session_duration	0.002069	-0.000845	
0.008153			
page_views	0.003101	-0.003670	
0.002657			
cart_abandonment_rate	-0.000313	0.004395	-
0.005104			
average_spent	0.008415	-0.004868	
0.009656			
user_income	0.000450	0.000561	-
0.000427			
product_return_rate	-0.001619	0.002045	
0.001719			
discount_percentage	1.000000	-0.001379	-
0.003327			
time_to_purchase	-0.001379	1.000000	
0.003150			
delivery_time	-0.003327	0.003150	
1.000000			
shipping_fee	-0.005994	-0.005498	-
0.005238			
seller_response_time	-0.007436	-0.001043	-
0.003248			
product_rating_variance	0.008739	0.004788	
0.004401			
review_sentiment_score	-0.008945	-0.003292	-
0.000429			
user_engagement_score	-0.000484	-0.003193	-
0.005483			
ad_click_rate	-0.002277	0.003012	-

0.004506		
product_popularity	-0.004310	-0.002184
0.004569		

	shipping_fee	seller_response_time \
price	0.001692	0.000019
review_count	0.003383	-0.000159
user_age	-0.000638	-0.003831
time_on_page	0.009305	-0.000462
session_duration	-0.000676	0.007093
page_views	0.002222	0.000090
cart_abandonment_rate	0.001536	0.003947
average_spent	0.003372	0.000699
user_income	0.003680	-0.003094
product_return_rate	0.004030	0.001394
discount_percentage	-0.005994	-0.007436
time_to_purchase	-0.005498	-0.001043
delivery_time	-0.005238	-0.003248
shipping_fee	1.000000	0.006113
seller_response_time	0.006113	1.000000
product_rating_variance	-0.002822	-0.004816
review_sentiment_score	0.000610	-0.000052
user_engagement_score	0.001722	0.002366
ad_click_rate	-0.000414	-0.004610
product_popularity	-0.001354	0.000244

	product_rating_variance	
review_sentiment_score \		
price	0.011372	-
0.009115		
review_count	-0.004260	
0.002156		
user_age	0.002916	-
0.003348		
time_on_page	0.003512	
0.000330		
session_duration	-0.000166	
0.001741		
page_views	0.000418	-
0.002711		
cart_abandonment_rate	0.007186	-
0.003193		
average_spent	-0.002859	
0.006323		
user_income	0.001301	
0.005753		
product_return_rate	0.003918	
0.003248		
discount_percentage	0.008739	-

0.008945		
time_to_purchase	0.004788	-
0.003292		
delivery_time	0.004401	-
0.000429		
shipping_fee	-0.002822	
0.000610		
seller_response_time	-0.004816	-
0.000052		
product_rating_variance	1.000000	
0.005371		
review_sentiment_score	0.005371	
1.000000		
user_engagement_score	0.003763	
0.004033		
ad_click_rate	-0.001256	
0.003371		
product_popularity	0.002944	-
0.007109		


	user_engagement_score	ad_click_rate	\
price	-0.000325	0.004739	
review_count	0.006545	0.005577	
user_age	-0.001121	0.001147	
time_on_page	-0.000154	-0.000799	
session_duration	-0.000691	0.002220	
page_views	-0.003882	0.003758	
cart_abandonment_rate	0.003944	0.007245	
average_spent	-0.000966	-0.005269	
user_income	-0.000097	-0.004528	
product_return_rate	0.001266	0.001472	
discount_percentage	-0.000484	-0.002277	
time_to_purchase	-0.003193	0.003012	
delivery_time	-0.005483	-0.004506	
shipping_fee	0.001722	-0.000414	
seller_response_time	0.002366	-0.004610	
product_rating_variance	0.003763	-0.001256	
review_sentiment_score	0.004033	0.003371	
user_engagement_score	1.000000	0.004989	
ad_click_rate	0.004989	1.000000	
product_popularity	0.004979	0.005217	

	product_popularity
price	-0.001021
review_count	0.004671
user_age	0.002571
time_on_page	-0.004101
session_duration	-0.003204
page_views	-0.005096

cart_abandonment_rate	0.005538
average_spent	0.004336
user_income	-0.002315
product_return_rate	-0.002570
discount_percentage	-0.004310
time_to_purchase	-0.002184
delivery_time	0.004569
shipping_fee	-0.001354
seller_response_time	0.000244
product_rating_variance	0.002944
review_sentiment_score	-0.007109
user_engagement_score	0.004979
ad_click_rate	0.005217
product_popularity	1.000000

1.b) Heatmap

```
plt.figure(figsize=(20,20))  
sns.heatmap(corrmatrix,cmap='viridis',vmin=-1,vmax=1,annot=True)  
plt.show()
```



2) Feature Selection using Stats

2.a) ANNOVA Testing

```
import scipy.stats as stats
from scipy.stats import chi2_contingency

print("ANNOVA on Numerical vs Target Column")
print()
annovanumerical={}
```

```

numerical_to_keep=[]
numerical_can_be_removed=[]
for i in numerical:
    group1=df[df['purchase_history']==True][i]
    group2=df[df['purchase_history']==False][i]
    f_stat, p_value = stats.f_oneway(group1,group2)
    annovanumerical[i]=p_value
    if p_value<0.05:
        numerical_to_keep.append(i)
    elif p_value>0.05:
        numerical_can_be_removed.append(i)
print("column is important for prediction \n",numerical_to_keep)
print()
print("column is not important, can be removed \n",numerical_can_be_removed)
print()
print(annovanumerical)

```

ANNOVA on Numerical vs Target Column

column is important for prediction
 ['average_spent', 'user_income', 'time_to_purchase',
 'seller_response_time', 'product_popularity']

column is not important, can be removed
 ['price', 'review_count', 'user_age', 'time_on_page',
 'session_duration', 'page_views', 'cart_abandonment_rate',
 'product_return_rate', 'discount_percentage', 'delivery_time',
 'shipping_fee', 'product_rating_variance', 'review_sentiment_score',
 'user_engagement_score', 'ad_click_rate']

```

{'price': 0.8899465739986773, 'review_count': 0.5770768106977051,
'user_age': 0.7854081232478979, 'time_on_page': 0.5085345397926913,
'session_duration': 0.7497831381683305, 'page_views':
0.1449642634414101, 'cart_abandonment_rate': 0.07615900506164756,
'average_spent': 0.03134234160871035, 'user_income':
0.015616516056511025, 'product_return_rate': 0.5786243875252175,
'discount_percentage': 0.5028842277197112, 'time_to_purchase':
0.0009285761046717416, 'delivery_time': 0.715906387899155,
'shipping_fee': 0.12814595886083166, 'seller_response_time':
0.011110530302949694, 'product_rating_variance': 0.9687752069566598,
'review_sentiment_score': 0.5063007444466396, 'user_engagement_score':
0.43557617269905147, 'ad_click_rate': 0.478543850161665,
'product_popularity': 0.019631567972866765}

```

2.b) ChiSquare Testing

```

print("Chi Square Test on Categorical Data")
print()

```

```

chisquarecategorical={}
categorical_to_keep=[]
categorical_to_remove=[]
for i in categorical:
    contingency_table=pd.crosstab(df[i],df['purchase_history'])
    chi2, p_value, dof, expected =chi2_contingency(contingency_table)
    chisquarecategorical[i]=p_value
    if p_value<0.05:
        categorical_to_keep.append(i)

    elif p_value>0.05:
        categorical_to_remove.append(i)
print("column is important for prediction \n",categorical_to_keep)
print()
print("column is not important, can be removed \
n",categorical_to_remove)
print()
print(chisquarecategorical)

```

Chi Square Test on Categorical Data

```

column is important for prediction
[]

```

```

column is not important, can be removed
['category', 'user_gender', 'user_location', 'search_keywords',
'user_membership', 'user_browser', 'user_device', 'referral_source',
'user_education', 'user_marital_status', 'product_availability',
'stock_status', 'product_color', 'product_size', 'seller_location',
'time_of_day', 'day_of_week', 'season', 'payment_method']

```

```

{'category': 0.4953927713533365, 'user_gender': 0.5103166694448208,
'user_location': 0.6412943409626277, 'search_keywords':
0.11971165254357352, 'user_membership': 0.14738046643188624,
'user_browser': 0.8750665879204095, 'user_device': 0.6032563538695197,
'referral_source': 0.5704558268399278, 'user_education':
0.05188126377538571, 'user_marital_status': 0.5860140869575095,
'product_availability': 0.9171695966427554, 'stock_status':
0.5717932793932304, 'product_color': 0.5774925654273336,
'product_size': 0.736638141818841, 'seller_location':
0.6536341963997004, 'time_of_day': 0.6395547722074488, 'day_of_week':
0.42747766253353114, 'season': 0.9392576700028451, 'payment_method':
0.3067340087346109}

```

```

print("Chi Square Test on discrete Data")
print()

```

```

chisquarediscrete={}
discrete_to_keep=[]
discrete_to_remove=[]

```

```

for i in discrete:
    contingency_table=pd.crosstab(df[i],df['purchase_history'])
    chi2, p_value, dof, expected =chi2_contingency(contingency_table)
    chisquaredescrete[i]=p_value
    if p_value<0.05:
        discrete_to_keep.append(i)
    elif p_value>0.05:
        discrete_to_remove.append(i)
print("column is important for prediction \n",discrete_to_keep)
print()
print("column is not important, can be removed \n",discrete_to_remove)
print()
print(chisquaredescrete)

```

Chi Square Test on discrete Data

column is important for prediction
[]

column is not important, can be removed
['rating', 'add_to_cart_count', 'discount_applied', 'clicks_on_ads',
'wishlist_additions', 'is_top_seller', 'seller_rating', 'coupon_used',
'purchase_time_month']

```

{'rating': 0.5379206436964865, 'add_to_cart_count':
0.7324985410378212, 'discount_applied': 0.09075982734615025,
'clicks_on_ads': 0.6597429536690044, 'wishlist_additions':
0.372469260506758, 'is_top_seller': 0.5561143004611644,
'seller_rating': 0.550055705756747, 'coupon_used': 0.7631196614510456,
'purchase_time_month': 0.12279738470928081}

```

5) Observations:

total duplicate rows in dataset is = 0

total missing values rows in dataset is = 0

target column "purchase_history" is perfectly balanced

multicollinearity is not there in numerical data

boxplot - No Outliers Present

ANNOVA on Numerical Dependent Variable vs Target Categorical Column: 5 column are important for prediction 'average_spent' 'user_income', 'time_to_purchase', 'seller_response_time', 'product_popularity We can try building a model using these columns only'##### Chi Square Test on Categorical Variable vs Target Categorical Column0 5 column are important for prediction

II] Encoding, Train-Test-Split, & Standardization

1) Label Encoding

```
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
for i in categorical:
    x[i]=le.fit_transform(x[i])
```

```
y=pd.DataFrame(np.where(y==True,1,0))
```

```
for i in discrete:
    if x[i].dtype=='bool':
        x[i] = x[i].astype(int)
```

```
x.head()
```

	category	price	rating	review_count	user_age	user_gender	\
0	0	842.23	2	155	24	2	
1	0	253.76	3	331	43	2	
2	0	483.65	2	236	64	0	
3	4	459.37	2	227	34	0	
4	4	150.11	2	214	51	0	

	user_location	time_on_page	add_to_cart_count
search_keywords	...	\	
0	2	13.86	6
4	...		
1	1	13.03	3
1	...		
2	0	3.75	7
4	...		
3	2	6.01	0
0	...		
4	2	6.89	9
2	...		

	review_sentiment_score	user_engagement_score	ad_click_rate
time_of_day	\		
0	-0.28	0.68	0.04
3			
1	0.28	0.11	0.89
2			
2	0.23	0.35	0.99
1			
3	0.93	0.73	0.16
0			
4	0.11	0.26	0.17
3			

	day_of_week	season	payment_method	coupon_used
product_popularity \				
0	4	2	2	0
0.54				
1	2	2	2	0
0.77				
2	5	0	2	0
0.14				
3	5	1	1	0
0.18				
4	6	1	3	0
0.66				

	purchase_time_month
0	8
1	7
2	9
3	9
4	9

[5 rows x 48 columns]

2) Train Test Split

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=42)
x_train2,x_test2,y_train2,y_test2=train_test_split(x,y,test_size=0.3,random_state=42)
for i in [x_train,x_test,y_train,y_test]:
    print(i.shape)

(42000, 48)
(18000, 48)
(42000, 1)
(18000, 1)
```

3) Standard Scaling

```
x_train.shape
(42000, 48)

x_test.shape
(18000, 48)

from sklearn.preprocessing import StandardScaler
ss=StandardScaler()
for i in numerical:
```

```

x_train[i]=ss.fit_transform(x_train[[i]])
x_test[i]=ss.transform(x_test[[i]])

x_train.shape
(42000, 48)

x_test.shape
(18000, 48)

from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report, precision_score, recall_score, f1_score

def evaluation_matrix(a,b):
    print(f"accuracy = {accuracy_score(a,b)}")
    print(f"precision = {precision_score(a,b)}")
    print(f"recall = {recall_score(a,b)}")
    print(f"f1 score = {f1_score(a,b)}")
    print("Confusion Matrix:\n", confusion_matrix(a,b))

```

4) Model Evaluation Metrics Function

```

model_metrics=pd.DataFrame(columns=['Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])
model_metrics

Empty DataFrame
Columns: [Model, Accuracy, Precision, Recall, F1 Score]
Index: []

def model_evaluation_metrics(model_name,ytrain,ypredict):
    accuracy=accuracy_score(ytrain,ypredict)
    precision=precision_score(ytrain,ypredict)
    recall=recall_score(ytrain,ypredict)
    flscore=f1_score(ytrain,ypredict)

metrics=pd.DataFrame([[model_name,accuracy,precision,recall,flscore]],
columns=['Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])
global model_metrics
model_metrics=pd.concat([model_metrics,metrics],ignore_index=True)

```

III] Model Building

1) Logistic Regression

```

from sklearn.linear_model import LogisticRegression
lr=LogisticRegression()

```



```

lr.fit(x_train[numerical_to_keep],y_train)
y_train_predict=lr.predict(x_train[numerical_to_keep])
y_test_predict=lr.predict(x_test[numerical_to_keep])

evaluation_matrix(y_train,y_train_predict)

accuracy = 0.5131190476190476
precision = 0.5135257191889084
recall = 0.5402413530976815
f1 score = 0.5265448820356092
Confusion Matrix:
[[10180 10772]
 [ 9677 11371]]

evaluation_matrix(y_test,y_test_predict)

accuracy = 0.5043888888888889
precision = 0.5063076433796247
recall = 0.5283185840707965
f1 score = 0.5170789801331673
Confusion Matrix:
[[4303 4657]
 [4264 4776]]

model_evaluation_metrics("Logistic Regression -
Train",y_train,y_train_predict)
model_evaluation_metrics("Logistic Regression -
Test",y_test,y_test_predict)

model_metrics

```

	Model	Accuracy	Precision	Recall	F1
Score					
0	Logistic Regression - Train	0.513119	0.513526	0.540241	0.526545
1	Logistic Regression - Test	0.504389	0.506308	0.528319	0.517079

2) Logistic Regression With Standard Scaling on Categorical Columns

```

from sklearn.preprocessing import StandardScaler
ss=StandardScaler()
for i in x_train2.columns:
    x_train2[i]=ss.fit_transform(x_train2[[i]])
    x_test2[i]=ss.transform(x_test2[[i]])

lr2=LogisticRegression()
lr2.fit(x_train2[numerical_to_keep],y_train2)

```

```
y_train2_predict=lr.predict(x_train2[numerical_to_keep])
y_test2_predict=lr.predict(x_test2[numerical_to_keep])
```

```
evaluation_matrix(y_train2,y_train2_predict)
```

```
accuracy = 0.5131190476190476
precision = 0.5135257191889084
recall = 0.5402413530976815
f1 score = 0.5265448820356092
```

```
Confusion Matrix:
```

```
[[10180 10772]
 [ 9677 11371]]
```

```
evaluation_matrix(y_test2,y_test2_predict)
```

```
accuracy = 0.5043888888888889
precision = 0.5063076433796247
recall = 0.5283185840707965
f1 score = 0.5170789801331673
```

```
Confusion Matrix:
```

```
[[4303 4657]
 [4264 4776]]
```

```
model_evaluation_metrics("Logistic Regression with SS -
Train",y_train2,y_train2_predict)
```

```
model_evaluation_metrics("Logistic Regression with SS -
Test",y_test2,y_test2_predict)
```

```
model_metrics
```

	Model	Accuracy	Precision	Recall
0	Logistic Regression - Train	0.513119	0.513526	0.540241
1	Logistic Regression - Test	0.504389	0.506308	0.528319
2	Logistic Regression with SS - Train	0.513119	0.513526	0.540241
3	Logistic Regression with SS - Test	0.504389	0.506308	0.528319

	F1 Score
0	0.526545
1	0.517079
2	0.526545
3	0.517079

3) Stats Model

```
import statsmodels.api as sm
x_train3 = sm.add_constant(x_train)
logistic=sm.Logit(y_train,x_train3)
model=logistic.fit()
print(model.summary())
```

```
Optimization terminated successfully.
      Current function value: 0.692129
      Iterations 3
```

Logit Regression Results

```
=====
Dep. Variable:                  0    No. Observations:
42000
Model:                        Logit    Df Residuals:
41951
Method:                       MLE     Df Model:
48
Date:                Sun, 16 Feb 2025    Pseudo R-squ.:
0.001465
Time:                21:14:56    Log-Likelihood:
-29069.
converged:                True    LL-Null:
-29112.
Covariance Type:            nonrobust    LLR p-value:
0.0007366
=====
```

```
=====
[0.025    0.975]
-----
coef      std err          z      P>|z|
-----
const                0.0828    0.075    1.108    0.268
-0.064    0.229
category             -0.0088    0.007   -1.273    0.203
-0.022    0.005
price               -0.0039    0.010   -0.401    0.688
-0.023    0.015
rating              -0.0011    0.007   -0.156    0.876
-0.015    0.012
review_count        -0.0035    0.010   -0.362    0.717
-0.023    0.016
user_age            -0.0062    0.010   -0.631    0.528
-0.025    0.013
user_gender          0.0155    0.012    1.297    0.195
-0.008    0.039
user_location       -0.0120    0.012   -1.003    0.316
```

-0.036	0.011				
time_on_page		-0.0112	0.010	-1.145	0.252
-0.030	0.008				
add_to_cart_count		0.0040	0.003	1.182	0.237
-0.003	0.011				
search_keywords		-0.0117	0.007	-1.682	0.093
-0.025	0.002				
discount_applied		-0.0385	0.020	-1.971	0.049
-0.077	-0.000				
user_membership		0.0087	0.009	1.002	0.317
-0.008	0.026				
user_browser		0.0043	0.009	0.489	0.625
-0.013	0.021				
user_device		-0.0057	0.012	-0.477	0.633
-0.029	0.018				
session_duration		0.0021	0.010	0.215	0.830
-0.017	0.021				
clicks_on_ads		-0.0021	0.002	-1.213	0.225
-0.005	0.001				
page_views		0.0083	0.010	0.849	0.396
-0.011	0.027				
referral_source		-0.0060	0.009	-0.684	0.494
-0.023	0.011				
wishlist_additions		5.303e-05	0.002	0.031	0.975
-0.003	0.003				
cart_abandonment_rate		0.0063	0.010	0.648	0.517
-0.013	0.025				
average_spent		-0.0262	0.010	-2.676	0.007
-0.045	-0.007				
user_income		0.0204	0.010	2.087	0.037
0.001	0.040				
user_education		0.0172	0.009	1.966	0.049
5.15e-05	0.034				
user_marital_status		0.0091	0.009	1.046	0.296
-0.008	0.026				
product_availability		0.0104	0.012	0.866	0.387
-0.013	0.034				
stock_status		0.0097	0.012	0.809	0.418
-0.014	0.033				
product_return_rate		-0.0031	0.010	-0.319	0.750
-0.022	0.016				
product_color		0.0053	0.007	0.769	0.442
-0.008	0.019				
product_size		0.0123	0.012	1.026	0.305
-0.011	0.036				
is_top_seller		-0.0088	0.020	-0.451	0.652
-0.047	0.029				
discount_percentage		0.0023	0.010	0.237	0.812
-0.017	0.021				

time_to_purchase	0.0356	0.010	3.637	0.000
0.016 0.055				
delivery_time	-0.0104	0.010	-1.065	0.287
-0.030 0.009				
shipping_fee	-0.0218	0.010	-2.227	0.026
-0.041 -0.003				
seller_rating	-0.0094	0.007	-1.357	0.175
-0.023 0.004				
seller_response_time	0.0255	0.010	2.613	0.009
0.006 0.045				
seller_location	0.0078	0.012	0.654	0.513
-0.016 0.031				
product_rating_variance	0.0036	0.010	0.367	0.714
-0.016 0.023				
review_sentiment_score	-0.0060	0.010	-0.610	0.542
-0.025 0.013				
user_engagement_score	0.0109	0.010	1.112	0.266
-0.008 0.030				
ad_click_rate	0.0042	0.010	0.434	0.664
-0.015 0.023				
time_of_day	-0.0130	0.009	-1.481	0.139
-0.030 0.004				
day_of_week	-0.0070	0.005	-1.432	0.152
-0.017 0.003				
season	-0.0014	0.009	-0.162	0.871
-0.019 0.016				
payment_method	-0.0151	0.009	-1.726	0.084
-0.032 0.002				
coupon_used	0.0058	0.020	0.298	0.766
-0.032 0.044				
product_popularity	0.0271	0.010	2.770	0.006
0.008 0.046				
purchase_time_month	-0.0028	0.003	-0.985	0.325
-0.008 0.003				

```
=====
=====
```

```
# Predict probabilities
```

```
y_train_prob = model.predict(x_train3) # Predict on train set
```

```
y_test_prob = model.predict(sm.add_constant(x_test)) # Add constant
to test set before prediction
```

```
# Convert probabilities to class labels (threshold = 0.5)
```

```
y_train_pred = (y_train_prob > 0.5).astype(int)
```

```
y_test_pred = (y_test_prob > 0.5).astype(int)
```

```
evaluation_matrix(y_train,y_train_pred)
```

```
accuracy = 0.5175952380952381
```

```
precision = 0.5179787088225888
```

```

recall = 0.5386259977194983
f1 score = 0.5281006172120647
Confusion Matrix:
[[10402 10550]
 [ 9711 11337]]

```

```
evaluation_matrix(y_test,y_test_pred)
```

```

accuracy = 0.5028888888888889
precision = 0.5049601035152038
recall = 0.5180309734513274
f1 score = 0.5114120345091188
Confusion Matrix:
[[4369 4591]
 [4357 4683]]

```

```

model_evaluation_metrics("Stats Model (LR) -
Train",y_train,y_train_predict)
model_evaluation_metrics("Stats Model (LR) -
Test",y_test,y_test_predict)
model_metrics

```

	Model	Accuracy	Precision	Recall
0	Logistic Regression - Train	0.513119	0.513526	0.540241
1	Logistic Regression - Test	0.504389	0.506308	0.528319
2	Logistic Regression with SS - Train	0.513119	0.513526	0.540241
3	Logistic Regression with SS - Test	0.504389	0.506308	0.528319
4	Stats Model (LR) - Train	0.513119	0.513526	0.540241
5	Stats Model (LR) - Test	0.504389	0.506308	0.528319

```

F1 Score
0 0.526545
1 0.517079
2 0.526545
3 0.517079
4 0.526545
5 0.517079

```

```
## Get the p-values
```

```
# p_values = model.pvalues
```

```
## Define the significance level (e.g., 0.05)
```

```
# alpha = 0.05
```

```
# # Get the columns with p-values less than alpha
# significant_columns = p_values[p_values <= alpha].index

# print(significant_columns)
```

4) Decision Tree

```
from sklearn.tree import DecisionTreeClassifier
dt=DecisionTreeClassifier(criterion='log_loss',max_depth=4)

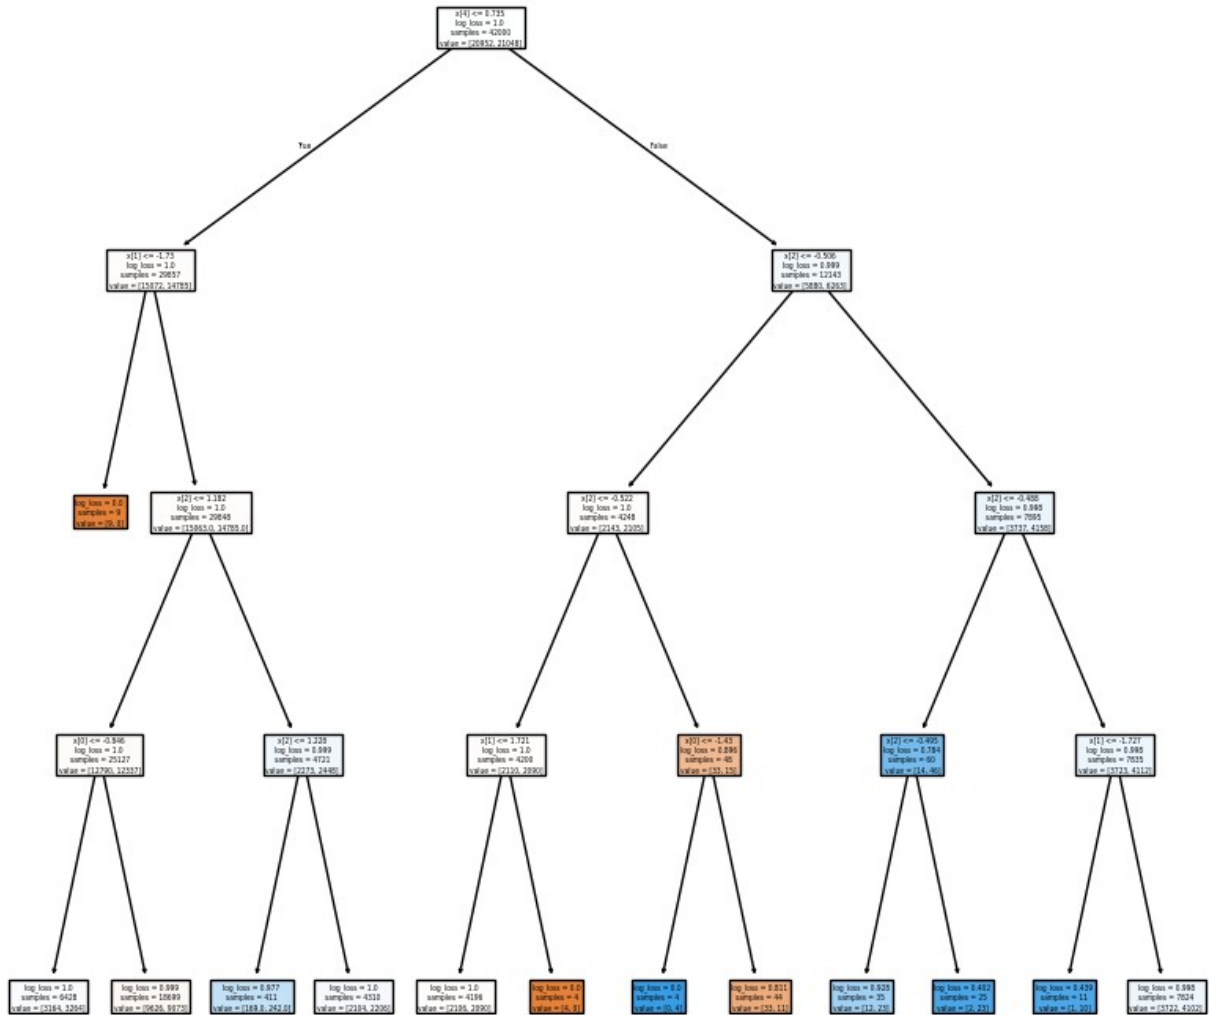
dt.fit(x_train[numerical_to_keep],y_train)

DecisionTreeClassifier(criterion='log_loss', max_depth=4)
```

4.a) DT Tree Diagram

```
from sklearn.tree import plot_tree
plt.figure(figsize=(10,10))
plot_tree(dt,filled=True)
plt.title("DT Tree Diagram")
plt.show()
```

DT Tree Diagram



```
y_train_predict=lr.predict(x_train[numerical_to_keep])
evaluation_matrix(y_train,y_train_predict)
```

```
accuracy = 0.5131190476190476
precision = 0.5135257191889084
recall = 0.5402413530976815
f1 score = 0.5265448820356092
Confusion Matrix:
[[10180 10772]
 [ 9677 11371]]
```



```

y_test_predict=lr.predict(x_test[numerical_to_keep])
evaluation_matrix(y_test,y_test_predict)

accuracy = 0.5043888888888889
precision = 0.5063076433796247
recall = 0.5283185840707965
f1 score = 0.5170789801331673
Confusion Matrix:
[[4303 4657]
 [4264 4776]]

model_evaluation_metrics("DecisionTree -
Train",y_train,y_train_predict)
model_evaluation_metrics("DecisionTree - Test",y_test,y_test_predict)
model_metrics

```

	Model	Accuracy	Precision	Recall
0	Logistic Regression - Train	0.513119	0.513526	0.540241
1	Logistic Regression - Test	0.504389	0.506308	0.528319
2	Logistic Regression with SS - Train	0.513119	0.513526	0.540241
3	Logistic Regression with SS - Test	0.504389	0.506308	0.528319
4	Stats Model (LR) - Train	0.513119	0.513526	0.540241
5	Stats Model (LR) - Test	0.504389	0.506308	0.528319
6	DecisionTree - Train	0.513119	0.513526	0.540241
7	DecisionTree - Test	0.504389	0.506308	0.528319

```

F1 Score
0 0.526545
1 0.517079
2 0.526545
3 0.517079
4 0.526545
5 0.517079
6 0.526545
7 0.517079

```

5) Random Forest

```

from sklearn.ensemble import RandomForestClassifier

i=40
rfc=RandomForestClassifier(n_estimators=i,max_depth=4,random_state=42)

```

```

rfc.fit(x_train,y_train)
y_train_predict=rfc.predict(x_train)
y_test_predict=rfc.predict(x_test)
print()
print(f"n_estimator = {i}")
print(f"train_accuracy={accuracy_score(y_train,y_train_predict)},
test_accuracy={accuracy_score(y_test,y_test_predict)}")

n_estimator = 40
train_accuracy=0.5601666666666667, test_accuracy=0.5037777777777778

model_evaluation_metrics("Random Forest -
Train",y_train,y_train_predict)
model_evaluation_metrics("Random Forest - Test",y_test,y_test_predict)
model_metrics

```

	Model	Accuracy	Precision	Recall
0	Logistic Regression - Train	0.513119	0.513526	0.540241
1	Logistic Regression - Test	0.504389	0.506308	0.528319
2	Logistic Regression with SS - Train	0.513119	0.513526	0.540241
3	Logistic Regression with SS - Test	0.504389	0.506308	0.528319
4	Stats Model (LR) - Train	0.513119	0.513526	0.540241
5	Stats Model (LR) - Test	0.504389	0.506308	0.528319
6	DecisionTree - Train	0.513119	0.513526	0.540241
7	DecisionTree - Test	0.504389	0.506308	0.528319
8	Random Forest - Train	0.560167	0.554401	0.623385
9	Random Forest - Test	0.503778	0.505365	0.562721

```

F1 Score
0 0.526545
1 0.517079
2 0.526545
3 0.517079
4 0.526545
5 0.517079
6 0.526545
7 0.517079
8 0.586872
9 0.532503

```

5.a) Feature Importance

```
rfc.feature_importances_
```

```
array([0.00514383, 0.0219449 , 0.00305131, 0.03706279, 0.02525855,
        0.01459486, 0.00212649, 0.03665709, 0.0065069 , 0.0062489 ,
        0.00159372, 0.00837013, 0.00221248, 0.00408166, 0.03408216,
        0.01703656, 0.0219626 , 0.00187657, 0.03152289, 0.02445949,
        0.06159432, 0.05973389, 0.00633194, 0.00126244, 0.00744943,
        0.00182547, 0.01785313, 0.01209383, 0.00579391, 0.00582219,
        0.0477129 , 0.09691251, 0.03756426, 0.03618254, 0.00475532,
        0.05316106, 0.00230782, 0.04540263, 0.03572277, 0.04502118,
        0.02894014, 0.00756018, 0.01414072, 0.00401029, 0.00344366,
        0.00396357, 0.03187965, 0.0157644 ])
```

```
feature_importance=rfc.feature_importances_
```

```
features=x_train.columns
```

```
important_feature=pd.DataFrame({'Feature':features,'Importance':feature_importance})
```

```
important_feature
```

	Feature	Importance
0	category	0.005144
1	price	0.021945
2	rating	0.003051
3	review_count	0.037063
4	user_age	0.025259
5	user_gender	0.014595
6	user_location	0.002126
7	time_on_page	0.036657
8	add_to_cart_count	0.006507
9	search_keywords	0.006249
10	discount_applied	0.001594
11	user_membership	0.008370
12	user_browser	0.002212
13	user_device	0.004082
14	session_duration	0.034082
15	clicks_on_ads	0.017037
16	page_views	0.021963
17	referral_source	0.001877
18	wishlist_additions	0.031523
19	cart_abandonment_rate	0.024459
20	average_spent	0.061594
21	user_income	0.059734
22	user_education	0.006332
23	user_marital_status	0.001262
24	product_availability	0.007449
25	stock_status	0.001825
26	product_return_rate	0.017853
27	product_color	0.012094
28	product_size	0.005794

29	is_top_seller	0.005822
30	discount_percentage	0.047713
31	time_to_purchase	0.096913
32	delivery_time	0.037564
33	shipping_fee	0.036183
34	seller_rating	0.004755
35	seller_response_time	0.053161
36	seller_location	0.002308
37	product_rating_variance	0.045403
38	review_sentiment_score	0.035723
39	user_engagement_score	0.045021
40	ad_click_rate	0.028940
41	time_of_day	0.007560
42	day_of_week	0.014141
43	season	0.004010
44	payment_method	0.003444
45	coupon_used	0.003964
46	product_popularity	0.031880
47	purchase_time_month	0.015764

```
important_feature[important_feature['Importance']>=0.1]
```

Empty DataFrame

Columns: [Feature, Importance]

Index: []

```
(rfc.feature_importances_>=0.1).sum()
```

0

```
(rfc.feature_importances_>=0.07).sum()
```

1

```
important_feature[important_feature['Importance']>=0.07]
```

	Feature	Importance
31	time_to_purchase	0.096913

```
(rfc.feature_importances_>=0.05).sum()
```

4

```
important_feature[important_feature['Importance']>=0.05]
```

	Feature	Importance
20	average_spent	0.061594
21	user_income	0.059734
31	time_to_purchase	0.096913
35	seller_response_time	0.053161

```
(rfc.feature_importances_>=0.03).sum()
```

```

15
important_feature[important_feature['Importance']>=0.03].Feature
3          review_count
7          time_on_page
14         session_duration
18        wishlist_additions
20         average_spent
21         user_income
30        discount_percentage
31         time_to_purchase
32         delivery_time
33         shipping_fee
35        seller_response_time
37    product_rating_variance
38        review_sentiment_score
39        user_engagement_score
46        product_popularity
Name: Feature, dtype: object

```

6) Gradient Boosting

```

from sklearn.ensemble import GradientBoostingClassifier

gbc=GradientBoostingClassifier(learning_rate=0.001,
n_estimators=150,max_depth=4)
gbc.fit(x_train,y_train)

GradientBoostingClassifier(learning_rate=0.001, max_depth=4,
n_estimators=150)

y_train_predict=gbc.predict(x_train)
evaluation_matrix(y_train,y_train_predict)

accuracy = 0.5262619047619047
precision = 0.5192237031098641
recall = 0.7385024705435196
f1 score = 0.609747965087771
Confusion Matrix:
[[ 6559 14393]
 [ 5504 15544]]

y_test_predict=gbc.predict(x_test)
evaluation_matrix(y_test,y_test_predict)

accuracy = 0.5017777777777778
precision = 0.5027985074626866
recall = 0.7154867256637168
f1 score = 0.5905770635500365
Confusion Matrix:

```

```
[[2564 6396]
 [2572 6468]]
```

```
model_evaluation_metrics("Gradient Boosting -
Train",y_train,y_train_predict)
model_evaluation_metrics("Gradient Boosting -
Test",y_test,y_test_predict)
model_metrics
```

	Model	Accuracy	Precision	Recall
\				
0	Logistic Regression - Train	0.513119	0.513526	0.540241
1	Logistic Regression - Test	0.504389	0.506308	0.528319
2	Logistic Regression with SS - Train	0.513119	0.513526	0.540241
3	Logistic Regression with SS - Test	0.504389	0.506308	0.528319
4	Stats Model (LR) - Train	0.513119	0.513526	0.540241
5	Stats Model (LR) - Test	0.504389	0.506308	0.528319
6	DecisionTree - Train	0.513119	0.513526	0.540241
7	DecisionTree - Test	0.504389	0.506308	0.528319
8	Random Forest - Train	0.560167	0.554401	0.623385
9	Random Forest - Test	0.503778	0.505365	0.562721
10	Gradient Boosting - Train	0.526262	0.519224	0.738502
11	Gradient Boosting - Test	0.501778	0.502799	0.715487

	F1 Score
0	0.526545
1	0.517079
2	0.526545
3	0.517079
4	0.526545
5	0.517079
6	0.526545
7	0.517079
8	0.586872
9	0.532503
10	0.609748
11	0.590577

7) Adaboost

```
from sklearn.ensemble import AdaBoostClassifier

abc=AdaBoostClassifier(n_estimators=100,learning_rate=0.1,)
abc.fit(x_train,y_train)

AdaBoostClassifier(learning_rate=0.1, n_estimators=100)

y_train_predict=abc.predict(x_train)
evaluation_matrix(y_train,y_train_predict)

accuracy = 0.5242142857142857
precision = 0.5244928936111495
recall = 0.5417616875712656
f1 score = 0.5329874500455725
Confusion Matrix:
[[10614 10338]
 [ 9645 11403]]

y_test_predict=abc.predict(x_test)
evaluation_matrix(y_test,y_test_predict)

accuracy = 0.5043333333333333
precision = 0.5063250428816467
recall = 0.5224557522123894
f1 score = 0.51426393728223
Confusion Matrix:
[[4355 4605]
 [4317 4723]]

my_lr=LogisticRegression()
abc=AdaBoostClassifier(estimator=my_lr,n_estimators=40,learning_rate=0.1,)
abc.fit(x_train,y_train)

AdaBoostClassifier(estimator=LogisticRegression(), learning_rate=0.1,
                    n_estimators=40)

y_train_predict=abc.predict(x_train)
evaluation_matrix(y_train,y_train_predict)

accuracy = 0.518095238095238
precision = 0.5181573033707865
recall = 0.5477480045610034
f1 score = 0.5325419187953254
Confusion Matrix:
[[10231 10721]
 [ 9519 11529]]

y_test_predict=abc.predict(x_test)
evaluation_matrix(y_test,y_test_predict)
```

```

accuracy = 0.5007777777777778
precision = 0.5028559339961921
recall = 0.5258849557522124
f1 score = 0.5141126851951985
Confusion Matrix:
[[4260 4700]
 [4286 4754]]

my_rfc=RandomForestClassifier(n_estimators=40,max_depth=4,random_state
=42)
abc=AdaBoostClassifier(estimator=my_rfc,n_estimators=20,learning_rate=
0.1,)
abc.fit(x_train,y_train)

AdaBoostClassifier(estimator=RandomForestClassifier(max_depth=4,
                                                    n_estimators=40,
                                                    random_state=42),
                    learning_rate=0.1, n_estimators=20)

y_train_predict=abc.predict(x_train)
evaluation_matrix(y_train,y_train_predict)

accuracy = 0.594
precision = 0.5898542903399892
recall = 0.6231470923603193
f1 score = 0.6060438037149987
Confusion Matrix:
[[11832  9120]
 [ 7932 13116]]

y_test_predict=abc.predict(x_test)
evaluation_matrix(y_test,y_test_predict)

accuracy = 0.5049444444444444
precision = 0.5067901884408885
recall = 0.5325221238938053
f1 score = 0.5193376126004638
Confusion Matrix:
[[4275 4685]
 [4226 4814]]

model_evaluation_metrics("AdaBoost - Train",y_train,y_train_predict)
model_evaluation_metrics("AdaBoost - Test",y_test,y_test_predict)
model_metrics

```

	Model	Accuracy	Precision	Recall
0	Logistic Regression - Train	0.513119	0.513526	0.540241
1	Logistic Regression - Test	0.504389	0.506308	0.528319

2	Logistic Regression with SS - Train	0.513119	0.513526	0.540241
3	Logistic Regression with SS - Test	0.504389	0.506308	0.528319
4	Stats Model (LR) - Train	0.513119	0.513526	0.540241
5	Stats Model (LR) - Test	0.504389	0.506308	0.528319
6	DecisionTree - Train	0.513119	0.513526	0.540241
7	DecisionTree - Test	0.504389	0.506308	0.528319
8	Random Forest - Train	0.560167	0.554401	0.623385
9	Random Forest - Test	0.503778	0.505365	0.562721
10	Gradient Boosting - Train	0.526262	0.519224	0.738502
11	Gradient Boosting - Test	0.501778	0.502799	0.715487
12	AdaBoost - Train	0.594000	0.589854	0.623147
13	AdaBoost - Test	0.504944	0.506790	0.532522

	F1 Score
0	0.526545
1	0.517079
2	0.526545
3	0.517079
4	0.526545
5	0.517079
6	0.526545
7	0.517079
8	0.586872
9	0.532503
10	0.609748
11	0.590577
12	0.606044
13	0.519338

8) XGBoost

```
from xgboost import XGBClassifier

xgb=XGBClassifier(max_depth=2,max_leaf_nodes=3)
xgb.fit(x_train,y_train)

XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None,
```

```

early_stopping_rounds=None,
        enable_categorical=False, eval_metric=None,
feature_types=None,
        gamma=None, grow_policy=None, importance_type=None,
        interaction_constraints=None, learning_rate=None,
max_bin=None,
        max_cat_threshold=None, max_cat_to_onehot=None,
        max_delta_step=None, max_depth=2, max_leaf_nodes=3,
        max_leaves=None, min_child_weight=None, missing=nan,
        monotone_constraints=None, multi_strategy=None,
n_estimators=None,
        n_jobs=None, num_parallel_tree=None, ...)

```

```

y_train_predict=xgb.predict(x_train)
evaluation_matrix(y_train,y_train_predict)

```

```

accuracy = 0.5687857142857143
precision = 0.5693244582920266
recall = 0.572976054732041
f1 score = 0.5711444199758471
Confusion Matrix:
[[11829  9123]
 [ 8988 12060]]

```

```

y_test_predict=xgb.predict(x_test)
evaluation_matrix(y_test,y_test_predict)

```

```

accuracy = 0.5013888888888889
precision = 0.5035923510555985
recall = 0.5039823008849558
f1 score = 0.503787250511417
Confusion Matrix:
[[4469 4491]
 [4484 4556]]

```

```

model_evaluation_metrics("XGBoost - Train",y_train,y_train_predict)
model_evaluation_metrics("XGBoost - Test",y_test,y_test_predict)
model_metrics

```

	Model	Accuracy	Precision	Recall
0	Logistic Regression - Train	0.513119	0.513526	0.540241
1	Logistic Regression - Test	0.504389	0.506308	0.528319
2	Logistic Regression with SS - Train	0.513119	0.513526	0.540241
3	Logistic Regression with SS - Test	0.504389	0.506308	0.528319
4	Stats Model (LR) - Train	0.513119	0.513526	0.540241

5	Stats Model (LR) - Test	0.504389	0.506308	0.528319
6	DecisionTree - Train	0.513119	0.513526	0.540241
7	DecisionTree - Test	0.504389	0.506308	0.528319
8	Random Forest - Train	0.560167	0.554401	0.623385
9	Random Forest - Test	0.503778	0.505365	0.562721
10	Gradient Boosting - Train	0.526262	0.519224	0.738502
11	Gradient Boosting - Test	0.501778	0.502799	0.715487
12	AdaBoost - Train	0.594000	0.589854	0.623147
13	AdaBoost - Test	0.504944	0.506790	0.532522
14	XGBoost - Train	0.568786	0.569324	0.572976
15	XGBoost - Test	0.501389	0.503592	0.503982

	F1 Score
0	0.526545
1	0.517079
2	0.526545
3	0.517079
4	0.526545
5	0.517079
6	0.526545
7	0.517079
8	0.586872
9	0.532503
10	0.609748
11	0.590577
12	0.606044
13	0.519338
14	0.571144
15	0.503787

9) KNN

```
from sklearn.neighbors import KNeighborsClassifier

for i in range(5,30,5):
    knn=KNeighborsClassifier(n_neighbors=i)
    knn.fit(x_train,y_train)
    y_train_predict=knn.predict(x_train)
    print()
```

```

print(i,"Negihbors")
print("train accuracy=",accuracy_score(y_train,y_train_predict),";
test accuracy=",accuracy_score(y_test,y_test_predict))
y_test_predict=knn.predict(x_test)

```

5 Negihbors
train accuracy= 0.6852380952380952 ; test accuracy= 0.5013888888888889

10 Negihbors
train accuracy= 0.6198571428571429 ; test accuracy=
0.49994444444444447

15 Negihbors
train accuracy= 0.601 ; test accuracy= 0.4995555555555553

20 Negihbors
train accuracy= 0.5849047619047619 ; test accuracy= 0.5032777777777778

25 Negihbors
train accuracy= 0.5777142857142857 ; test accuracy= 0.5002777777777778

```

y_train_predict=knn.predict(x_train)
evaluation_matrix(y_train,y_train_predict)

```

```

accuracy = 0.5777142857142857
precision = 0.5811844298460633
recall = 0.5632364120106423
f1 score = 0.5720696810307387
Confusion Matrix:
[[12409  8543]
 [ 9193 11855]]

```

```

y_test_predict=knn.predict(x_test)
evaluation_matrix(y_test,y_test_predict)

```

```

accuracy = 0.5025
precision = 0.5048844960349386
recall = 0.48595132743362834
f1 score = 0.49523702158841104
Confusion Matrix:
[[4652 4308]
 [4647 4393]]

```

```

model_evaluation_metrics("KNN - Train",y_train,y_train_predict)
model_evaluation_metrics("KNN - Test",y_test,y_test_predict)
model_metrics

```

	Model	Accuracy	Precision	Recall
0	Logistic Regression - Train	0.513119	0.513526	0.540241

1	Logistic Regression - Test	0.504389	0.506308	0.528319
2	Logistic Regression with SS - Train	0.513119	0.513526	0.540241
3	Logistic Regression with SS - Test	0.504389	0.506308	0.528319
4	Stats Model (LR) - Train	0.513119	0.513526	0.540241
5	Stats Model (LR) - Test	0.504389	0.506308	0.528319
6	DecisionTree - Train	0.513119	0.513526	0.540241
7	DecisionTree - Test	0.504389	0.506308	0.528319
8	Random Forest - Train	0.560167	0.554401	0.623385
9	Random Forest - Test	0.503778	0.505365	0.562721
10	Gradient Boosting - Train	0.526262	0.519224	0.738502
11	Gradient Boosting - Test	0.501778	0.502799	0.715487
12	AdaBoost - Train	0.594000	0.589854	0.623147
13	AdaBoost - Test	0.504944	0.506790	0.532522
14	XGBoost - Train	0.568786	0.569324	0.572976
15	XGBoost - Test	0.501389	0.503592	0.503982
16	KNN - Train	0.577714	0.581184	0.563236
17	KNN - Test	0.502500	0.504884	0.485951
F1 Score				
0	0.526545			
1	0.517079			
2	0.526545			
3	0.517079			
4	0.526545			
5	0.517079			
6	0.526545			
7	0.517079			
8	0.586872			
9	0.532503			
10	0.609748			
11	0.590577			
12	0.606044			
13	0.519338			

```
14 0.571144
15 0.503787
16 0.572070
17 0.495237
```

10) SVM

```
from sklearn.svm import SVC
```

```
svc= SVC(kernel='rbf')
svc.fit(x_train,y_train)
```

```
SVC()
```

```
y_train_predict=svc.predict(x_train)
evaluation_matrix(y_train,y_train_predict)
```

```
accuracy = 0.5626666666666666
precision = 0.5626578135228655
recall = 0.5716932725199544
f1 score = 0.5671395579016826
Confusion Matrix:
[[11599  9353]
 [ 9015 12033]]
```

```
y_test_predict=svc.predict(x_test)
evaluation_matrix(y_test,y_test_predict)
```

```
accuracy = 0.5041111111111111
precision = 0.5063136907399203
recall = 0.5056415929203539
f1 score = 0.5059774186406907
Confusion Matrix:
[[4503 4457]
 [4469 4571]]
```

```
model_evaluation_metrics("SVM - Train",y_train,y_train_predict)
model_evaluation_metrics("SVM - Test",y_test,y_test_predict)
model_metrics
```

	Model	Accuracy	Precision	Recall
\				
0	Logistic Regression - Train	0.513119	0.513526	0.540241
1	Logistic Regression - Test	0.504389	0.506308	0.528319
2	Logistic Regression with SS - Train	0.513119	0.513526	0.540241
3	Logistic Regression with SS - Test	0.504389	0.506308	0.528319
4	Stats Model (LR) - Train	0.513119	0.513526	0.540241

5	Stats Model (LR) - Test	0.504389	0.506308	0.528319
6	DecisionTree - Train	0.513119	0.513526	0.540241
7	DecisionTree - Test	0.504389	0.506308	0.528319
8	Random Forest - Train	0.560167	0.554401	0.623385
9	Random Forest - Test	0.503778	0.505365	0.562721
10	Gradient Boosting - Train	0.526262	0.519224	0.738502
11	Gradient Boosting - Test	0.501778	0.502799	0.715487
12	AdaBoost - Train	0.594000	0.589854	0.623147
13	AdaBoost - Test	0.504944	0.506790	0.532522
14	XGBoost - Train	0.568786	0.569324	0.572976
15	XGBoost - Test	0.501389	0.503592	0.503982
16	KNN - Train	0.577714	0.581184	0.563236
17	KNN - Test	0.502500	0.504884	0.485951
18	SVM - Train	0.562667	0.562658	0.571693
19	SVM - Test	0.504111	0.506314	0.505642
	F1 Score			
0	0.526545			
1	0.517079			
2	0.526545			
3	0.517079			
4	0.526545			
5	0.517079			
6	0.526545			
7	0.517079			
8	0.586872			
9	0.532503			
10	0.609748			
11	0.590577			
12	0.606044			
13	0.519338			
14	0.571144			
15	0.503787			
16	0.572070			
17	0.495237			

```
18 0.567140
19 0.505977
```

```
model_evaluation_metrics("SVM - Train",y_train,y_train_predict)
model_evaluation_metrics("SVM - Test",y_test,y_test_predict)
model_metrics
```

	Model	Accuracy	Precision	Recall
\				
0	Logistic Regression - Train	0.513119	0.513526	0.540241
1	Logistic Regression - Test	0.504389	0.506308	0.528319
2	Logistic Regression with SS - Train	0.513119	0.513526	0.540241
3	Logistic Regression with SS - Test	0.504389	0.506308	0.528319
4	Stats Model (LR) - Train	0.513119	0.513526	0.540241
5	Stats Model (LR) - Test	0.504389	0.506308	0.528319
6	DecisionTree - Train	0.513119	0.513526	0.540241
7	DecisionTree - Test	0.504389	0.506308	0.528319
8	Random Forest - Train	0.560167	0.554401	0.623385
9	Random Forest - Test	0.503778	0.505365	0.562721
10	Gradient Boosting - Train	0.526262	0.519224	0.738502
11	Gradient Boosting - Test	0.501778	0.502799	0.715487
12	AdaBoost - Train	0.594000	0.589854	0.623147
13	AdaBoost - Test	0.504944	0.506790	0.532522
14	XGBoost - Train	0.568786	0.569324	0.572976
15	XGBoost - Test	0.501389	0.503592	0.503982
16	KNN - Train	0.577714	0.581184	0.563236
17	KNN - Test	0.502500	0.504884	0.485951
18	SVM - Train	0.562667	0.562658	0.571693
19	SVM - Test	0.504111	0.506314	0.505642
20	SVM - Train	0.562667	0.562658	0.571693
21	SVM - Test	0.504111	0.506314	0.505642

	F1 Score
0	0.526545
1	0.517079
2	0.526545
3	0.517079
4	0.526545
5	0.517079
6	0.526545
7	0.517079
8	0.586872
9	0.532503
10	0.609748
11	0.590577
12	0.606044
13	0.519338
14	0.571144
15	0.503787
16	0.572070
17	0.495237
18	0.567140
19	0.505977
20	0.567140
21	0.505977

```
file_name=r'C:\Users\venky\Downloads\model_metrics.csv'  
model_metrics.to_csv(file_name, index=False)
```

